



Calibration and Metacognition Metrics (lower is better)

Expected Calibration Error



Brier Score



Wrong@High-Confidence (0.9)

